

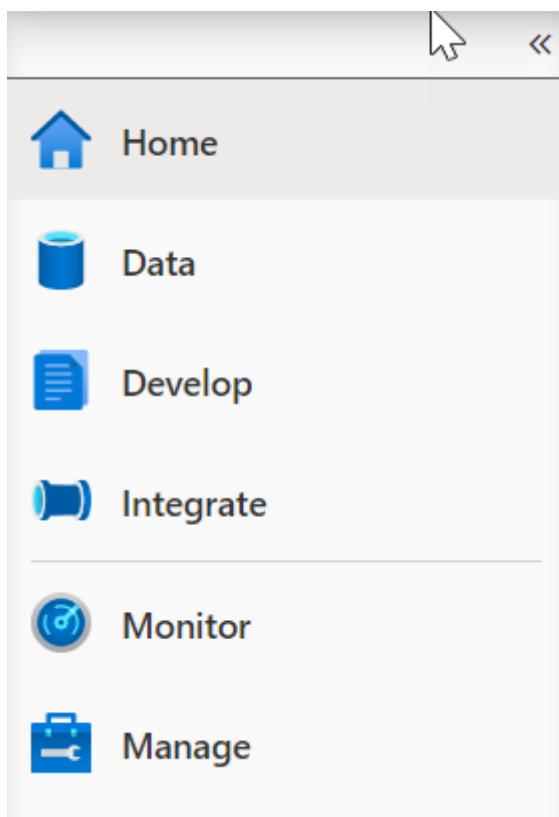
# Exercise - work with windowing functions

8 minutes

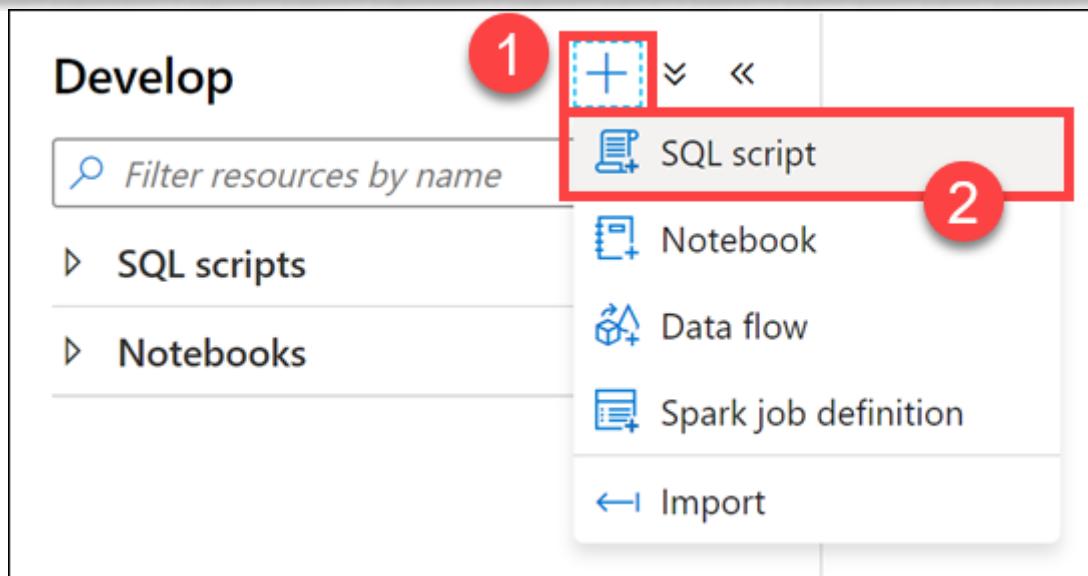
One of the key components of window functions is the OVER clause. This clause determines the partitioning and ordering of a rowset before the associated window function is applied. That is, the OVER clause defines a window or user-specified set of rows within a query result set. A window function then computes a value for each row in the window. You can use the OVER clause with functions to compute aggregated values such as moving averages, cumulative aggregates, running totals, or a top N per group results.

1. Open Synapse Studio (<https://web.azuresynapse.net/> ).

2. Select the Develop hub.



3. From the Develop menu, select the + button (1) and choose SQL Script (2) from the context menu.



4. In the toolbar menu, connect to the SQL Pool database to execute the query.



5. In the query window, replace the script with the following to use the OVER clause with data from the wwi\_security.Sale table:

```
SQL

SELECT
    ROW_NUMBER() OVER(PARTITION BY Region ORDER BY Quantity DESC) AS "Row Number",
    Product,
    Quantity,
    Region
FROM wwi_security.Sale
WHERE Quantity <> 0
ORDER BY Region;
```

6. Select Run from the toolbar menu to execute the SQL command.



When we use PARTITION BY with the OVER clause (1), we divide the query result set into partitions. The window function is applied to each partition separately and computation restarts for each partition.

```

1 SELECT
2     ROW_NUMBER() OVER(PARTITION BY Region ORDER BY Quantity DESC) AS "Row Number",
3     Product,
4     Quantity,
5     Region
6     FROM wwi_security.Sale
7     WHERE Quantity <> 0
8     ORDER BY Region;

```

Results    Messages

View    Table    Chart    Export results ▾

Search

Row Number	Product	Quantity	Region
16	Wheel	5	Far West
17	Video Games	5	Far West
18	Video Games	5	Far West
19	Wireless headphones	3	Far West
20	Wireless headphones	3	Far West
21	Key Chain	2	Far West
22	Key Chain	2	Far West
1	Sunglasses	52	South East
2	Beach hat	52	South East
	Beach hat	52	

The script we executed uses the OVER clause with ROW\_NUMBER function (1) to display a row number for each row within a partition. The partition in our case is the Region column. The ORDER BY clause (2) specified in the OVER clause orders the rows in each partition by the column Quantity. The ORDER BY clause in the SELECT statement determines the order in which the entire query result set is returned.

Scroll down in the results view until the Row Number count (3) starts over with a different region (4). Since the partition is set to Region, the ROW\_NUMBER resets when the region changes. Essentially, we have partitioned by region and have a result set identified by the number of rows in that region.

## Aggregate functions

Now use aggregate functions with our window by expanding on our query that uses the OVER clause.

1. In the query window, replace the script with the following to add aggregate functions:

SQL

```

SELECT
    ROW_NUMBER() OVER(PARTITION BY Region ORDER BY Quantity DESC) AS "Row

```

```

Number",
Product,
Quantity,
SUM(Quantity) OVER(PARTITION BY Region) AS Total,
AVG(Quantity) OVER(PARTITION BY Region) AS Avg,
COUNT(Quantity) OVER(PARTITION BY Region) AS Count,
MIN(Quantity) OVER(PARTITION BY Region) AS Min,
MAX(Quantity) OVER(PARTITION BY Region) AS Max,
Region
FROM wwi_security.Sale
WHERE Quantity <> 0
ORDER BY Region;

```

2. Select Run from the toolbar menu to execute the SQL command.



In our query, we added the SUM, AVG, COUNT, MIN, and MAX aggregate functions. Using the OVER clause is more efficient than using subqueries.

The screenshot shows a SQL query editor and its results. The query is:

```

5 SUM(Quantity) OVER(PARTITION BY Region) AS Total,
6 AVG(Quantity) OVER(PARTITION BY Region) AS Avg,
7 COUNT(Quantity) OVER(PARTITION BY Region) AS Count,
8 MIN(Quantity) OVER(PARTITION BY Region) AS Min,
9 MAX(Quantity) OVER(PARTITION BY Region) AS Max,
10 Region
11 FROM wwi_security.Sale
12 WHERE Quantity <> 0
13 ORDER BY Region;

```

The results table has columns: Row Number, Product, Quantity, Total, Avg, Count, Min, Max, and Region. The rows show data grouped by Region (Far West, South East). An arrow points from the highlighted 'OVER' clause in the query to the resulting table, indicating how the aggregate functions are applied across groups.

Row Number	Product	Quantity	Total	Avg	Count	Min	Max	Region
17	Video Games	3	320	14	22	2	32	Far West
18	Video Games	5	328	14	22	2	32	Far West
19	Wireless headphones	3	328	14	22	2	32	Far West
20	Wireless headphones	3	328	14	22	2	32	Far West
21	Key Chain	2	328	14	22	2	32	Far West
22	Key Chain	2	328	14	22	2	32	Far West
1	Sunglasses	52	962	32	30	5	52	South East
2	Beach hat	52	962	32	30	5	52	South East
3	Beach hat	52	962	32	30	5	52	South East
4	Sunglasses	52	962	32	30	5	52	South East

## Analytic functions

Analytic functions calculate an aggregate value based on a group of rows. Unlike aggregate functions, however, analytic functions can return multiple rows for each group. Use analytic functions to compute moving averages, running totals, percentages, or top-N results within a group.

Tailwind Traders has book sales data they import from their online store and wish to compute percentages of book downloads by category.

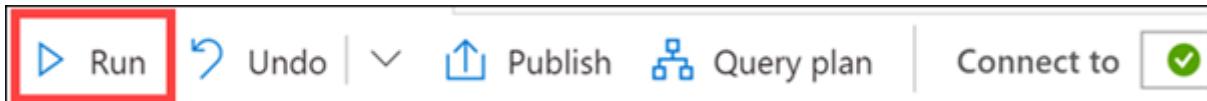
To do this, you decide to build window functions that use the PERCENTILE\_CONT and PERCENTILE\_DISC functions.

1. In the query window, replace the script with the following to add aggregate functions:

SQL

```
-- PERCENTILE_CONT, PERCENTILE_DISC
SELECT DISTINCT c.Category
,PERCENTILE_CONT(0.5) WITHIN GROUP (ORDER BY bc.Downloads)
    OVER (PARTITION BY Category) AS MedianCont
,PERCENTILE_DISC(0.5) WITHIN GROUP (ORDER BY bc.Downloads)
    OVER (PARTITION BY Category) AS MedianDisc
FROM dbo.Category AS c
INNER JOIN dbo.BookList AS bl
    ON bl.CategoryID = c.ID
INNER JOIN dbo.BookConsumption AS bc
    ON bc.BookID = bl.ID
ORDER BY Category
```

2. Select Run from the toolbar menu to execute the SQL command.



```

1 -- PERCENTILE_CONT, PERCENTILE_DISC
2 SELECT DISTINCT c.Category
3 ,PERCENTILE_CONT(0.5) WITHIN GROUP (ORDER BY bc.Downloads)
4 OVER (PARTITION BY Category) AS MedianCont
5 ,PERCENTILE_DISC(0.5) WITHIN GROUP (ORDER BY bc.Downloads)
6 OVER (PARTITION BY Category) AS MedianDisc
7 FROM dbo.Category AS c
8 INNER JOIN dbo.BookList AS bl
9 ON bl.CategoryID = c.ID
10 INNER JOIN dbo.BookConsumption AS bc
11 ON bc.BookID = bl.ID
12 ORDER BY Category

```

Results Messages

View Table Chart Export results ▾

Search

Category	MedianCont	MedianDisc
Favorites	304.75	231.833333333333
Fiction	1091.58333333333	935
Romance	78.583333333333	59.16666666666667
Sci-fi	54.75	44.33333333333333
Trending Now	187	153.333333333333

In this query, we use PERCENTILE\_CONT (1) and PERCENTILE\_DISC (2) to find the median number of downloads in each book category. These functions may not return the same value. PERCENTILE\_CONT interpolates the appropriate value, which may or may not exist in the data set, while PERCENTILE\_DISC always returns an actual value from the set. To explain further, PERCENTILE\_DISC computes a specific percentile for sorted values in an entire rowset or within a rowset's distinct partitions.

The 0.5 value passed to the percentile functions (1 & 2) computes the 50th percentile, or the median, of the downloads.

The WITHIN GROUP expression (3) specifies a list of values to sort and compute the percentile over. Only one ORDER BY expression is allowed, and the default sort order is ascending.

The OVER clause (4) divides the FROM clause's result set into partitions, in this case, Category. The percentile function is applied to these partitions.

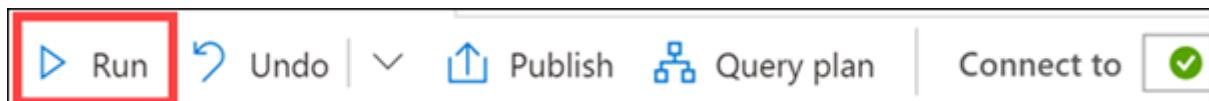
3. In the query window, replace the script with the following to use the LAG analytic function:

```
--LAG Function
SELECT ProductId,
       [Hour] AS SalesHour,
       TotalAmount AS CurrentSalesTotal,
       LAG(TotalAmount, 1,0) OVER (ORDER BY [Hour]) AS PreviousSalesTotal,
       TotalAmount - LAG(TotalAmount,1,0) OVER (ORDER BY [Hour]) AS Diff
FROM [wwi_perf].[Sale_Index]
WHERE ProductId = 3848 AND [Hour] BETWEEN 8 AND 20;
```

Tailwind Traders wants to compare the sales totals for a product over an hourly basis over time, showing the difference in value.

To accomplish this, you use the LAG analytic function. This function accesses data from a previous row in the same result set without the use of a self-join. LAG provides access to a row at a given physical offset that comes before the current row. We use this analytic function to compare values in the current row with values in a previous row.

4. Select Run from the toolbar menu to execute the SQL command.



```

1 --LAG Function
2 SELECT ProductId,
3     [Hour] AS SalesHour,
4     TotalAmount AS CurrentSalesTotal,
5     LAG(TotalAmount, 1,0) OVER (ORDER BY [Hour]) AS PreviousSalesTotal,
6     TotalAmount - LAG(TotalAmount,1,0) OVER (ORDER BY [Hour]) AS Diff
7 FROM [wwi_perf].[Sale_Index]
8 WHERE ProductId = 3848 AND [Hour] BETWEEN 8 AND 20;

```

Results    Messages

View    **Table**    Chart    Export results ▾

Search

ProductId	SalesHour	CurrentSalesTotal	PreviousSalesTotal	Diff
3848	8	27.71	0.00	27.71
3848	8	55.42	27.71	27.71
3848	8	27.71	55.42	-27.71
3848	8	110.84	27.71	83.13
3848	8	55.42	110.84	-55.42
3848	8	110.84	55.42	55.42
3848	8	110.84	110.84	0.00
3848	8	55.42	110.84	-55.42
3848	8	83.13	55.42	27.71
3848	8	27.71	83.13	-55.42

In this query, we use the LAG function (1) to return the difference in sales (2) for a specific product over peak sales hours (8-20). We also calculate the difference in sales from one row to the next (3). Notice that because there is no lag value available for the first row, the default of zero (0) is returned.

## ROWS clause

The ROWS and RANGE clauses further limit the rows within the partition by specifying start and end points within the partition. This is done by specifying a range of rows with respect to the current row either by logical association or physical association. Physical association is achieved by using the ROWS clause.

Tailwind Traders wants to find the books that have the fewest downloads by country, while displaying the total number of downloads for each book within each country in ascending order.

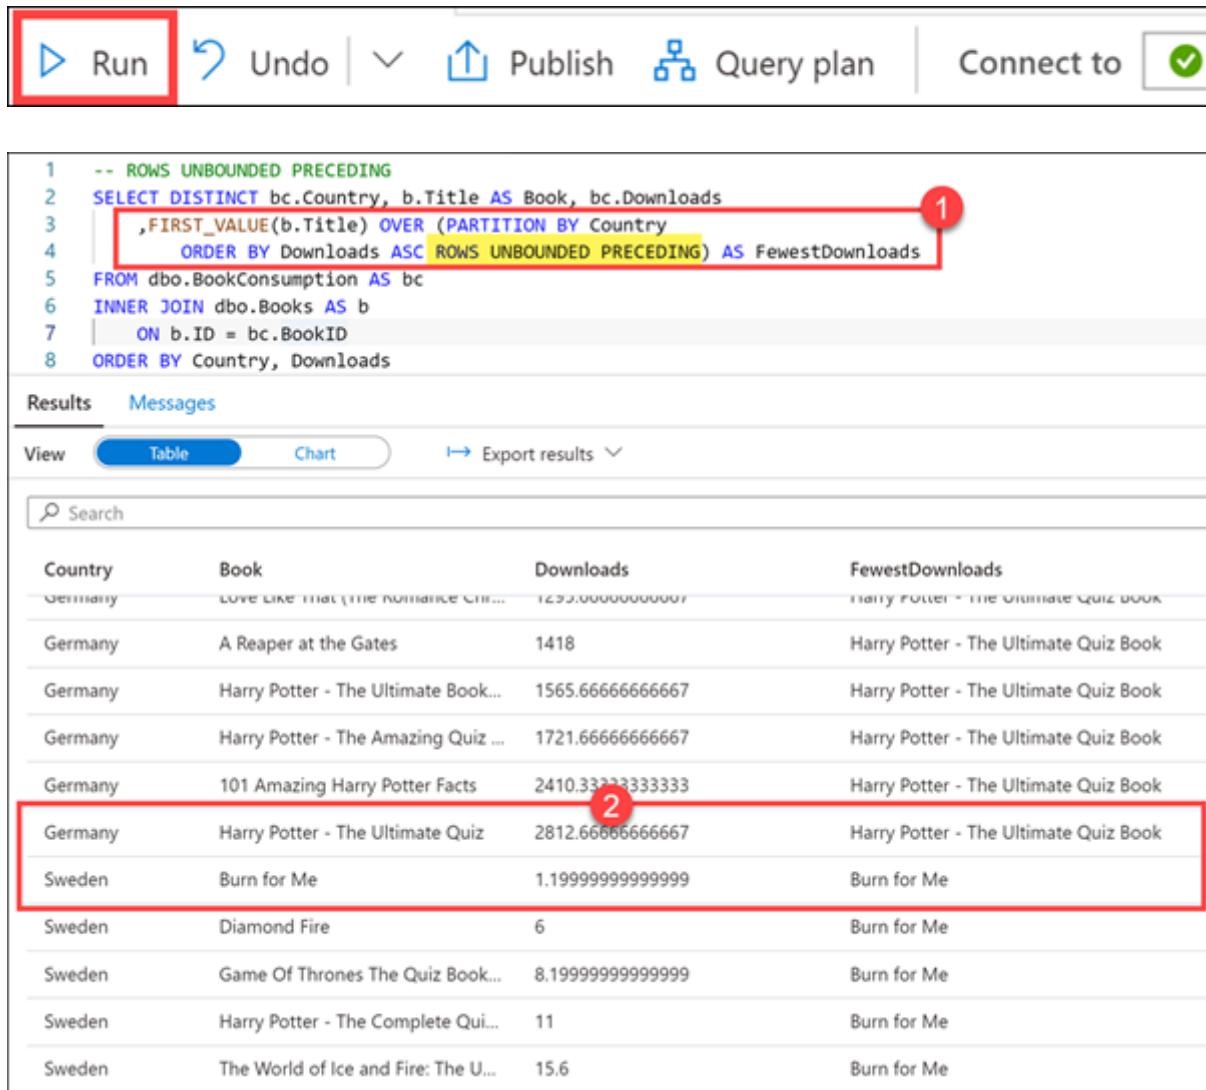
To achieve this, you use ROWS in combination with UNBOUNDED PRECEDING to limit the rows within the Country partition, specifying that the window start with the first row of the partition.

1. In the query window, replace the script with the following to add aggregate functions:

SQL

```
-- ROWS UNBOUNDED PRECEDING
SELECT DISTINCT bc.Country, b.Title AS Book, bc.Downloads
    ,FIRST_VALUE(b.Title) OVER (PARTITION BY Country
        ORDER BY Downloads ASC ROWS UNBOUNDED PRECEDING) AS Fewest-
Downloads
FROM dbo.BookConsumption AS bc
INNER JOIN dbo.Books AS b
    ON b.ID = bc.BookID
ORDER BY Country, Downloads
```

2. Select Run from the toolbar menu to execute the SQL command.



In this query, we use the FIRST\_VALUE analytic function to retrieve the book title with the fewest downloads, as indicated by the ROWS UNBOUNDED PRECEDING clause over the Country partition (1). The UNBOUNDED PRECEDING option set the window start to the first

row of the partition, giving us the title of the book with the fewest downloads for the country within the partition.

In the result set, we can scroll through the list that of books by country, sorted by number of downloads in ascending order.